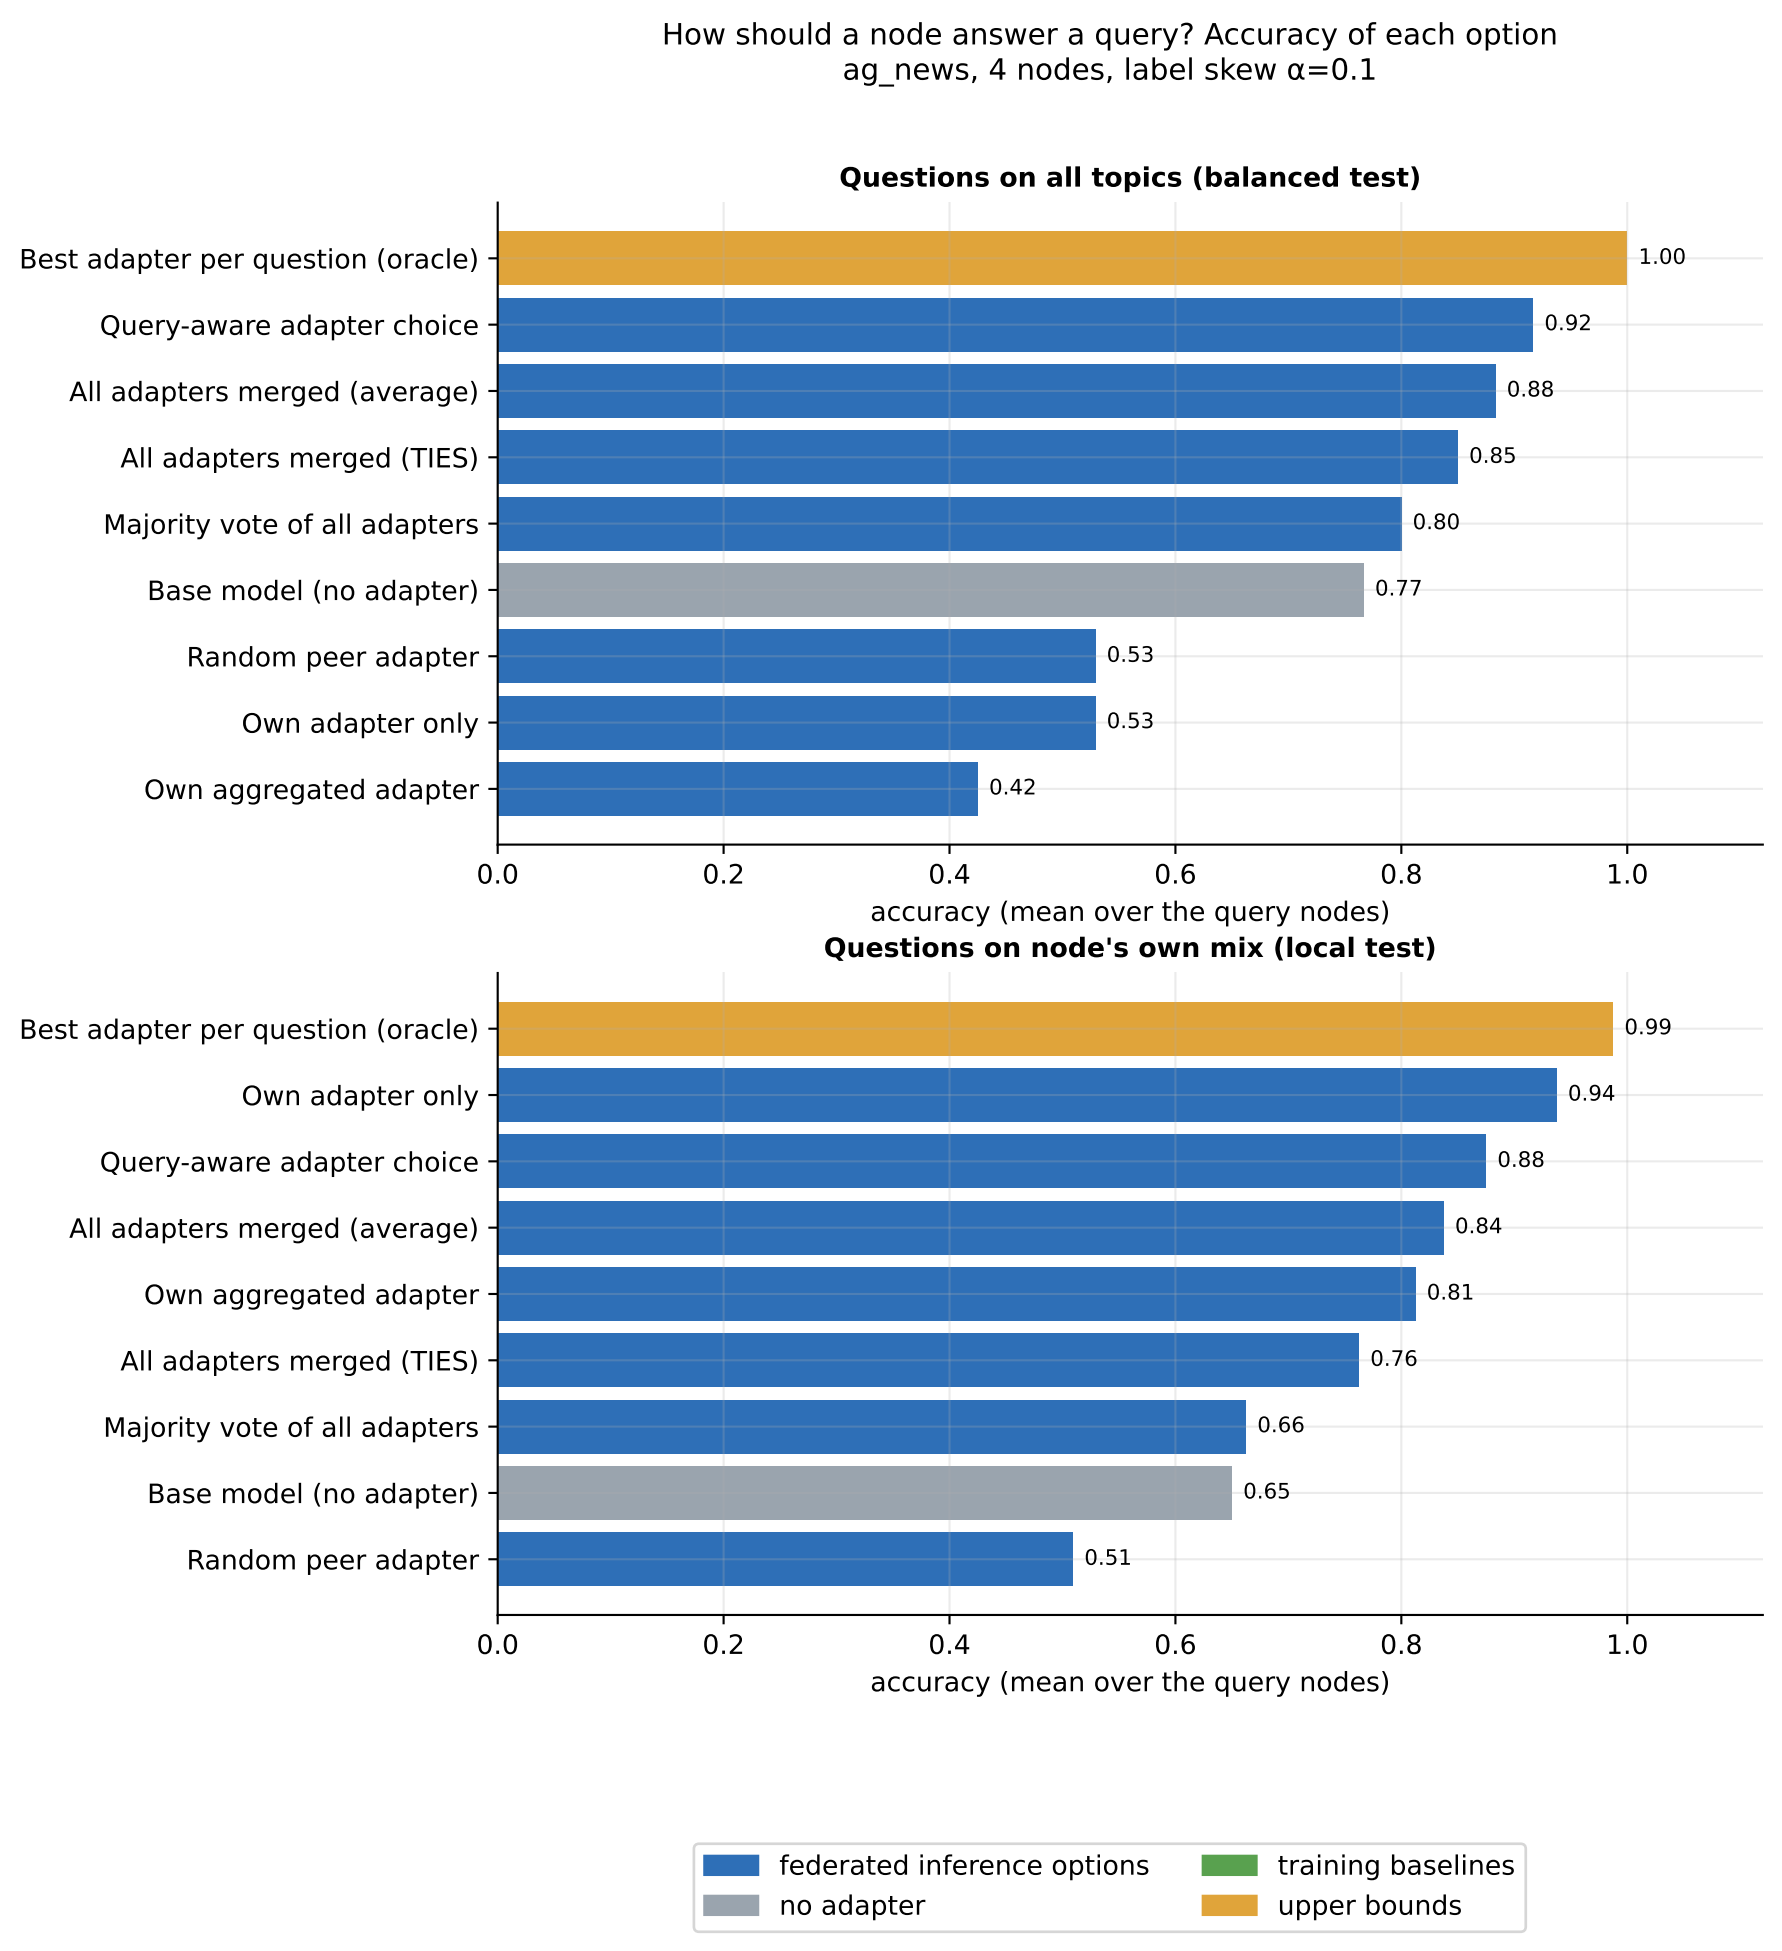
How should a node answer a query? Accuracy of each option
ag_news, 4 nodes, label skew α=0.1
Questions on all topics (balanced test)
Best adapter per question (oracle)
Query-aware adapter choice
All adapters merged (average)
All adapters merged (TIES)
Majority vote of all adapters
Base model (no adapter)
Random peer adapter
Own adapter only
Own aggregated adapter
1.00
0.92
0.88
0.85
0.80
0.77
0.53
0.53
0.42
0.0
0.2
0.4
0.6
0.8
1.0
accuracy (mean over the query nodes)
Questions on node's own mix (local test)
Best adapter per question (oracle)
Own adapter only
Query-aware adapter choice
All adapters merged (average)
Own aggregated adapter
All adapters merged (TIES)
Majority vote of all adapters
Base model (no adapter)
Random peer adapter
0.99
0.94
0.88
0.84
0.81
0.76
0.66
0.65
0.51
0.0
0.2
0.4
0.6
0.8
1.0
accuracy (mean over the query nodes)
federated inference options
no adapter
training baselines
upper bounds