Our federated inference vs a central multi-LoRA server and layer-split decentralised inference

	Ours: federated inference	Central multi-LoRA server (vLLM)	Layer-split decentralised (Petals-style)
TTFT, 256-token prompt (ms)	54.8	28.5	n/a
extra TTFT, first use of a peer adapter (ms)	436.5	0.0	n/a
TPOT (ms)	30.6	7.3	n/a
tokens/s, 16 requests on one machine	32.6	1,692.0	n/a
network latency per generated token (ms)	0.0	0.0	200.0
% answered when the preferred peer / the server is down	100.0	0.0	n/a
single point of failure	no	yes (the server)	no, if every layer range has a second server
needs a server	no	yes	no
new adapter known to peers after (ms, median)	14.5	n/a	n/a
where the adapters live	nowhere central (own + cached peers)	the server	split across servers