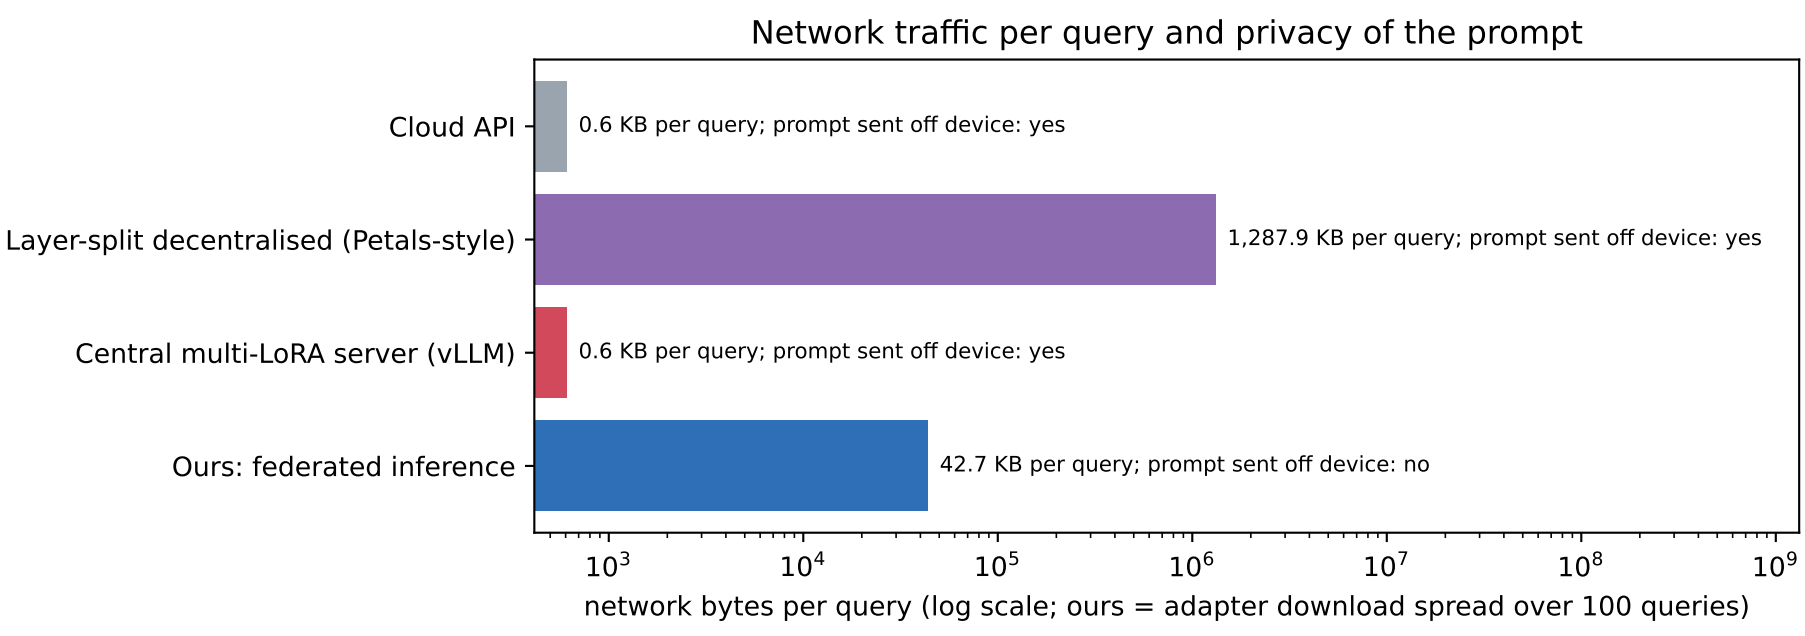
Network traffic per query and privacy of the prompt
Cloud API
0.6 KB per query; prompt sent off device: yes
Layer-split decentralised (Petals-style)
1,287.9 KB per query; prompt sent off device: yes
Central multi-LoRA server (vLLM)
0.6 KB per query; prompt sent off device: yes
Ours: federated inference
42.7 KB per query; prompt sent off device: no
10^3
10^4
10^5
10^6
10^7
10^8
10^9
network bytes per query (log scale; ours = adapter download spread over 100 queries)