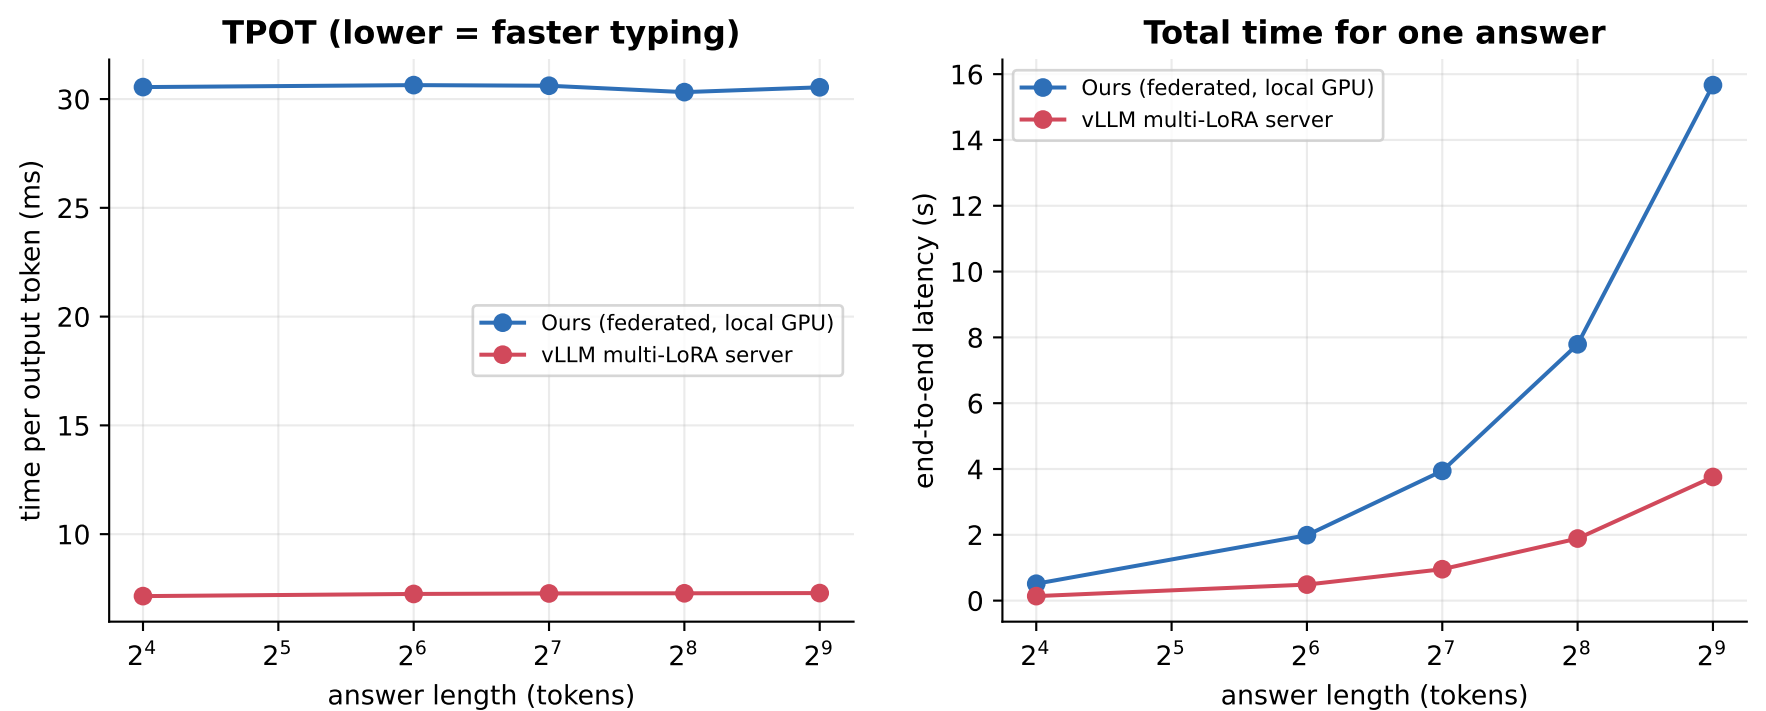
TPOT (lower = faster typing)
time per output token (ms)
30
25
20
15
10
Ours (federated, local GPU)
vLLM multi-LoRA server
2^4
2^5
2^6
2^7
2^8
2^9
answer length (tokens)
Total time for one answer
end-to-end latency (s)
16
14
12
10
8
6
4
2
0
Ours (federated, local GPU)
vLLM multi-LoRA server
2^4
2^5
2^6
2^7
2^8
2^9
answer length (tokens)