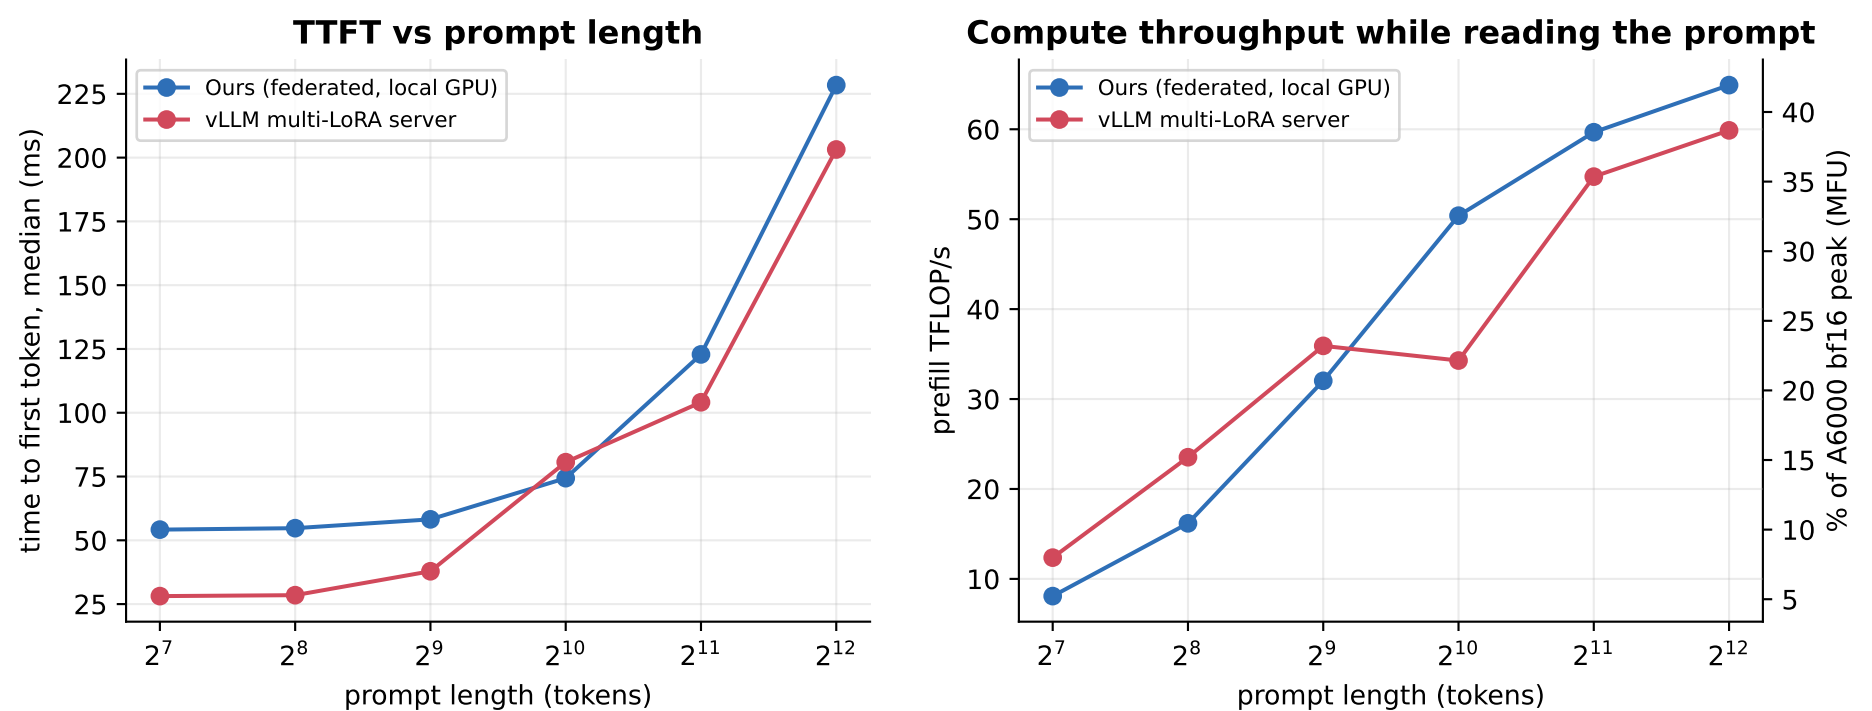
TTFT vs prompt length
Ours (federated, local GPU)
vLLM multi-LoRA server
time to first token, median (ms)
prompt length (tokens)
Compute throughput while reading the prompt
prefill TFLOP/s
% of A6000 bf16 peak (MFU)