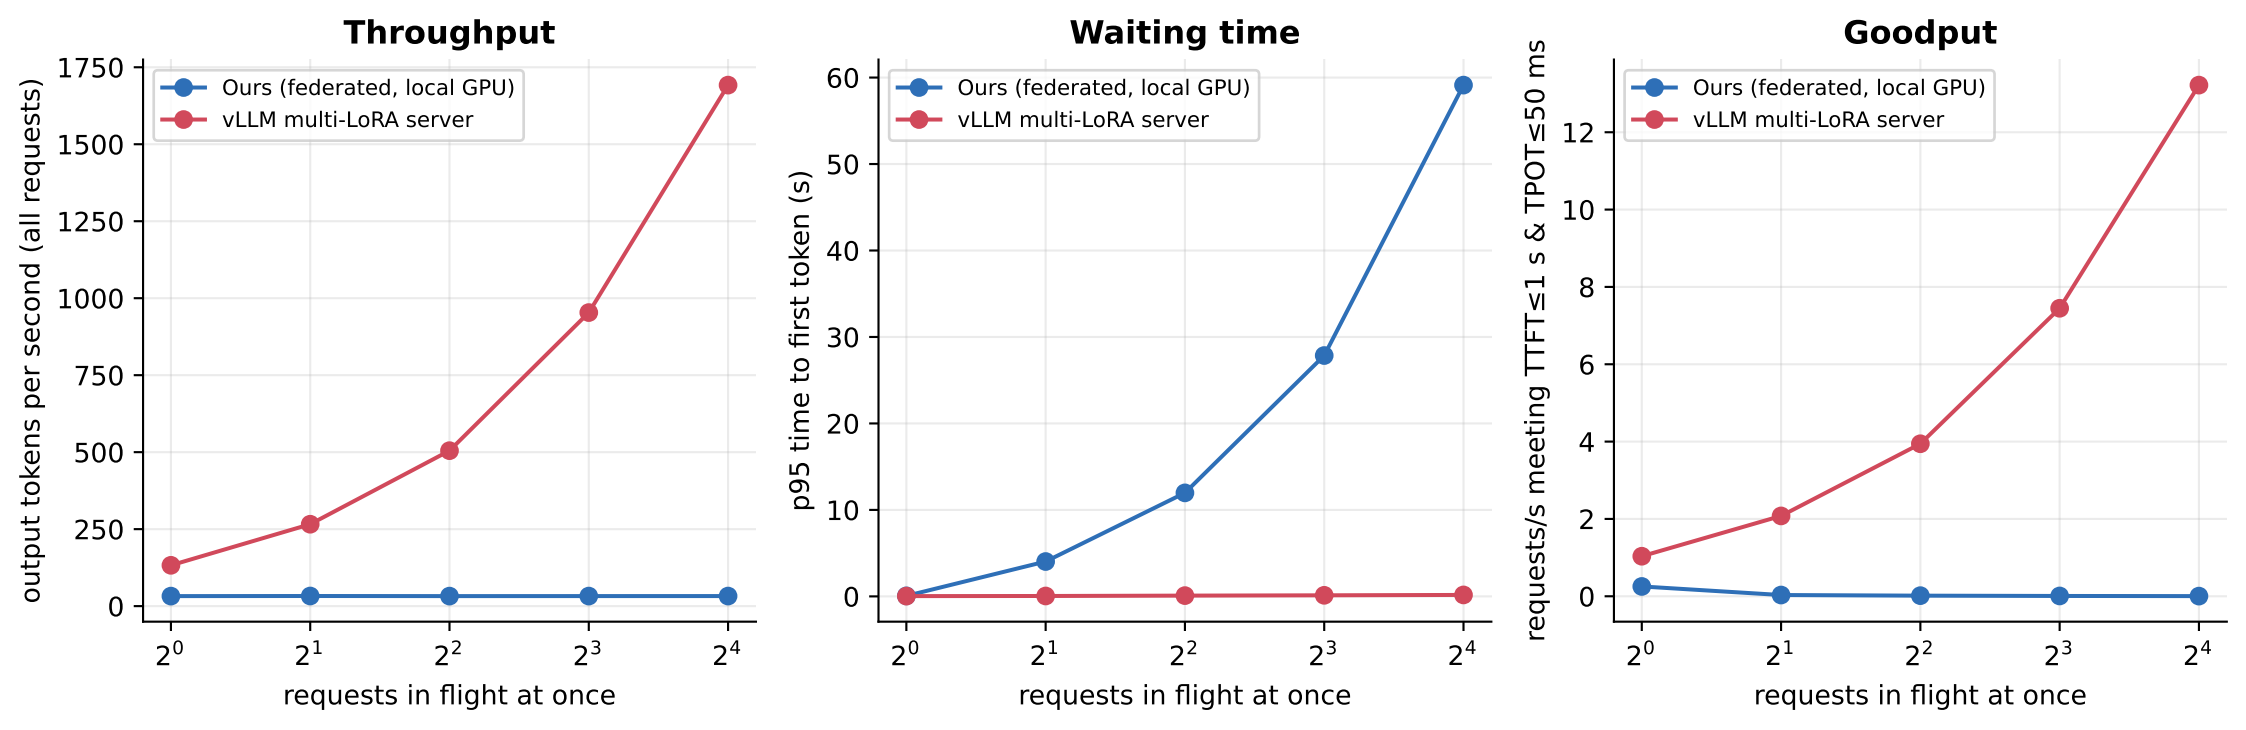
Throughput
output tokens per second (all requests)
1750
1500
1250
1000
750
500
250
0
Ours (federated, local GPU)
vLLM multi-LoRA server
2^0
2^1
2^2
2^3
2^4
requests in flight at once
Waiting time
p95 time to first token (s)
60
50
40
30
20
10
0
Ours (federated, local GPU)
vLLM multi-LoRA server
2^0
2^1
2^2
2^3
2^4
requests in flight at once
Goodput
requests/s meeting TTFT≤1 s & TPOT≤50 ms
12
10
8
6
4
2
0
Ours (federated, local GPU)
vLLM multi-LoRA server
2^0
2^1
2^2
2^3
2^4
requests in flight at once