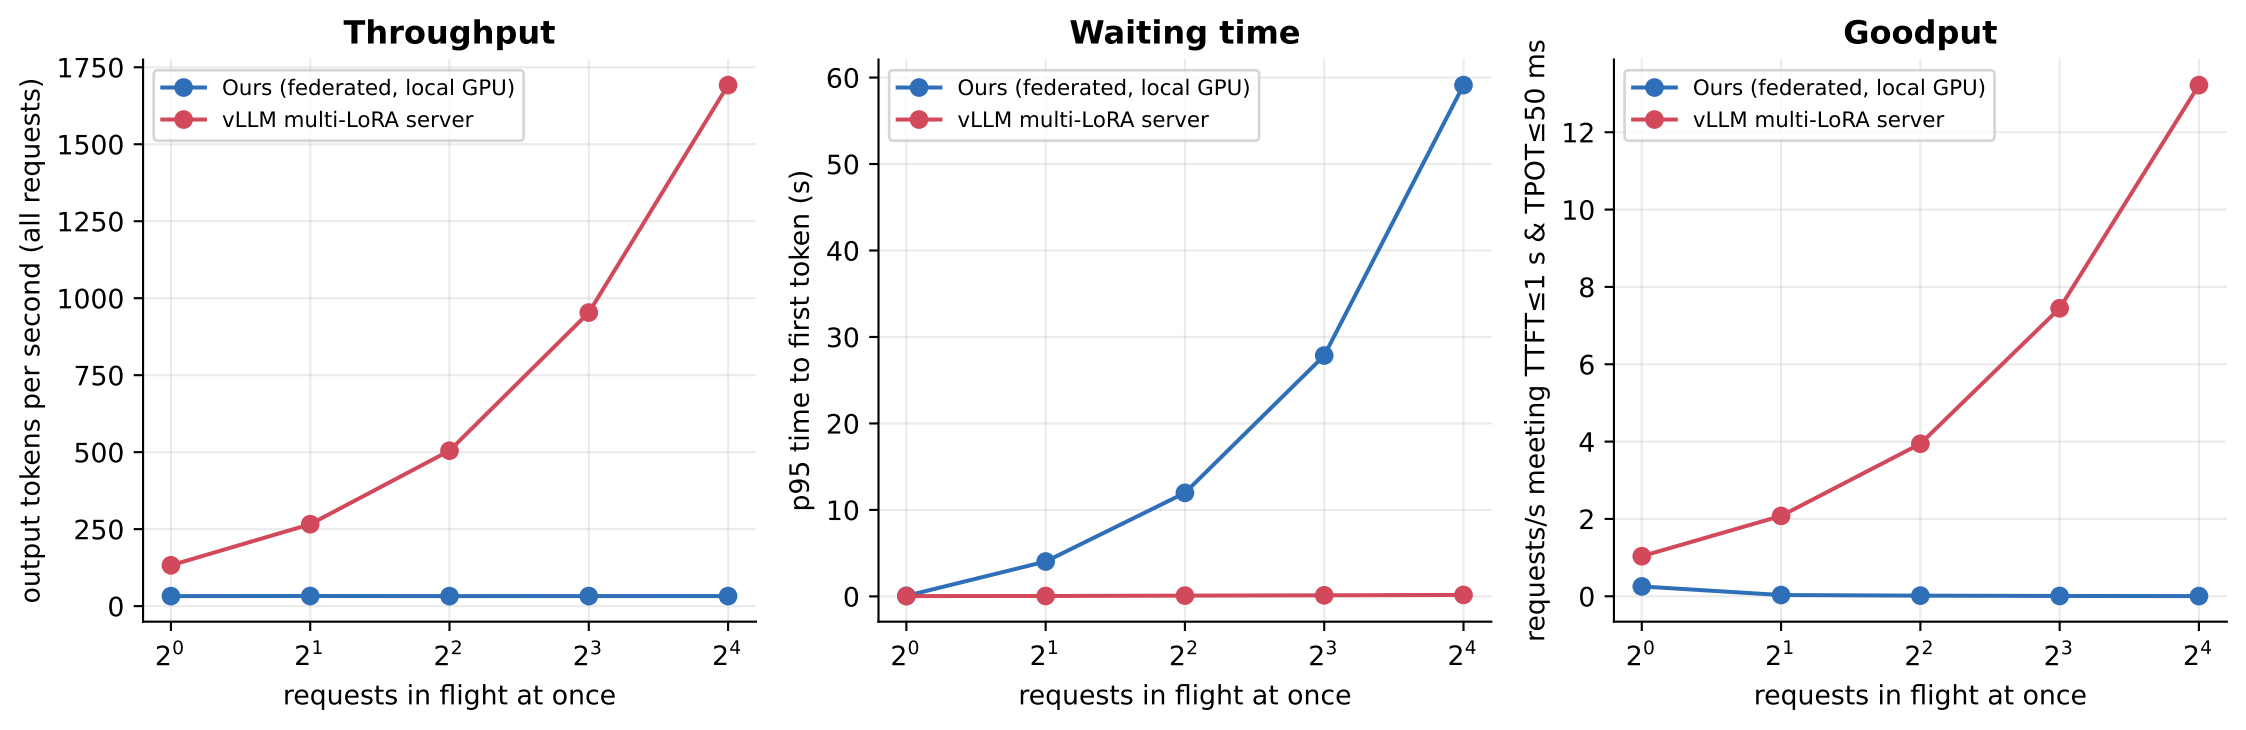
Throughput
output tokens per second (all requests)
1750
1500
1250
1000
750
500
250
0
Ours (federated, local GPU)
vLLM multi-LoRA server
$2^0$
$2^1$
$2^2$
$2^3$
$2^4$
requests in flight at once
Waiting time
p95 time to first token (s)
60
50
40
30
20
10
0
Ours (federated, local GPU)
vLLM multi-LoRA server
$2^0$
$2^1$
$2^2$
$2^3$
$2^4$
requests in flight at once
Goodput
requests/s meeting TTFT≤1 s & TPOT≤50 ms
12
10
8
6
4
2
0
Ours (federated, local GPU)
vLLM multi-LoRA server
$2^0$
$2^1$
$2^2$
$2^3$
$2^4$
requests in flight at once