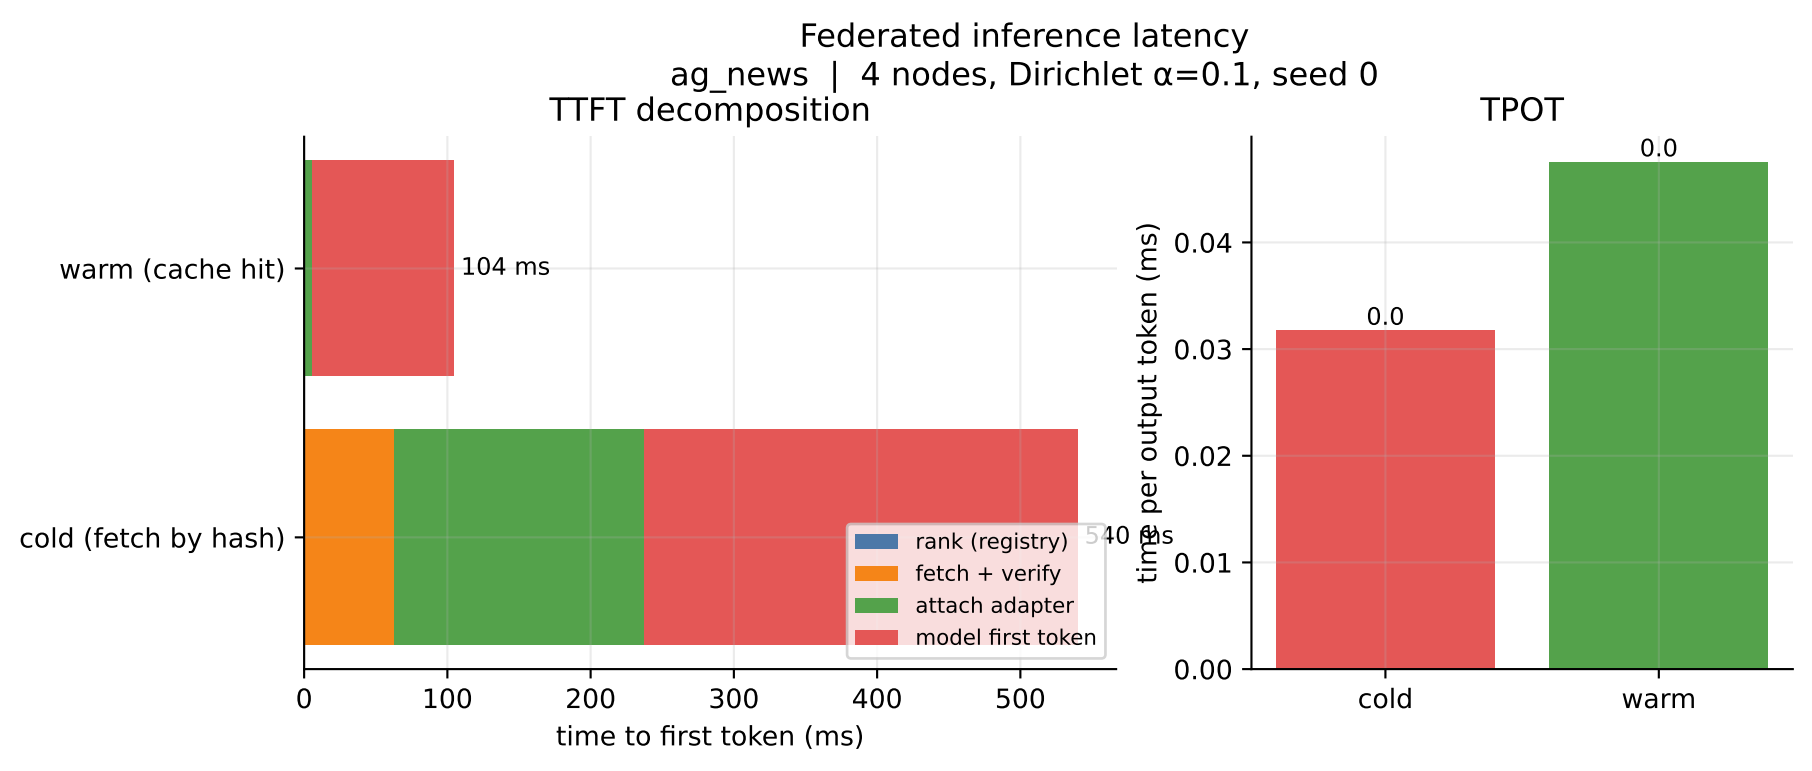
Federated inference latency
ag_news | 4 nodes, Dirichlet α=0.1, seed 0
TTFT decomposition
warm (cache hit)
cold (fetch by hash)
104 ms
540 ms
rank (registry)
fetch + verify
attach adapter
model first token
0
100
200
300
400
500
time to first token (ms)
TPOT
time per output token (ms)
0.00
0.01
0.02
0.03
0.04
0.0
0.0
cold
warm