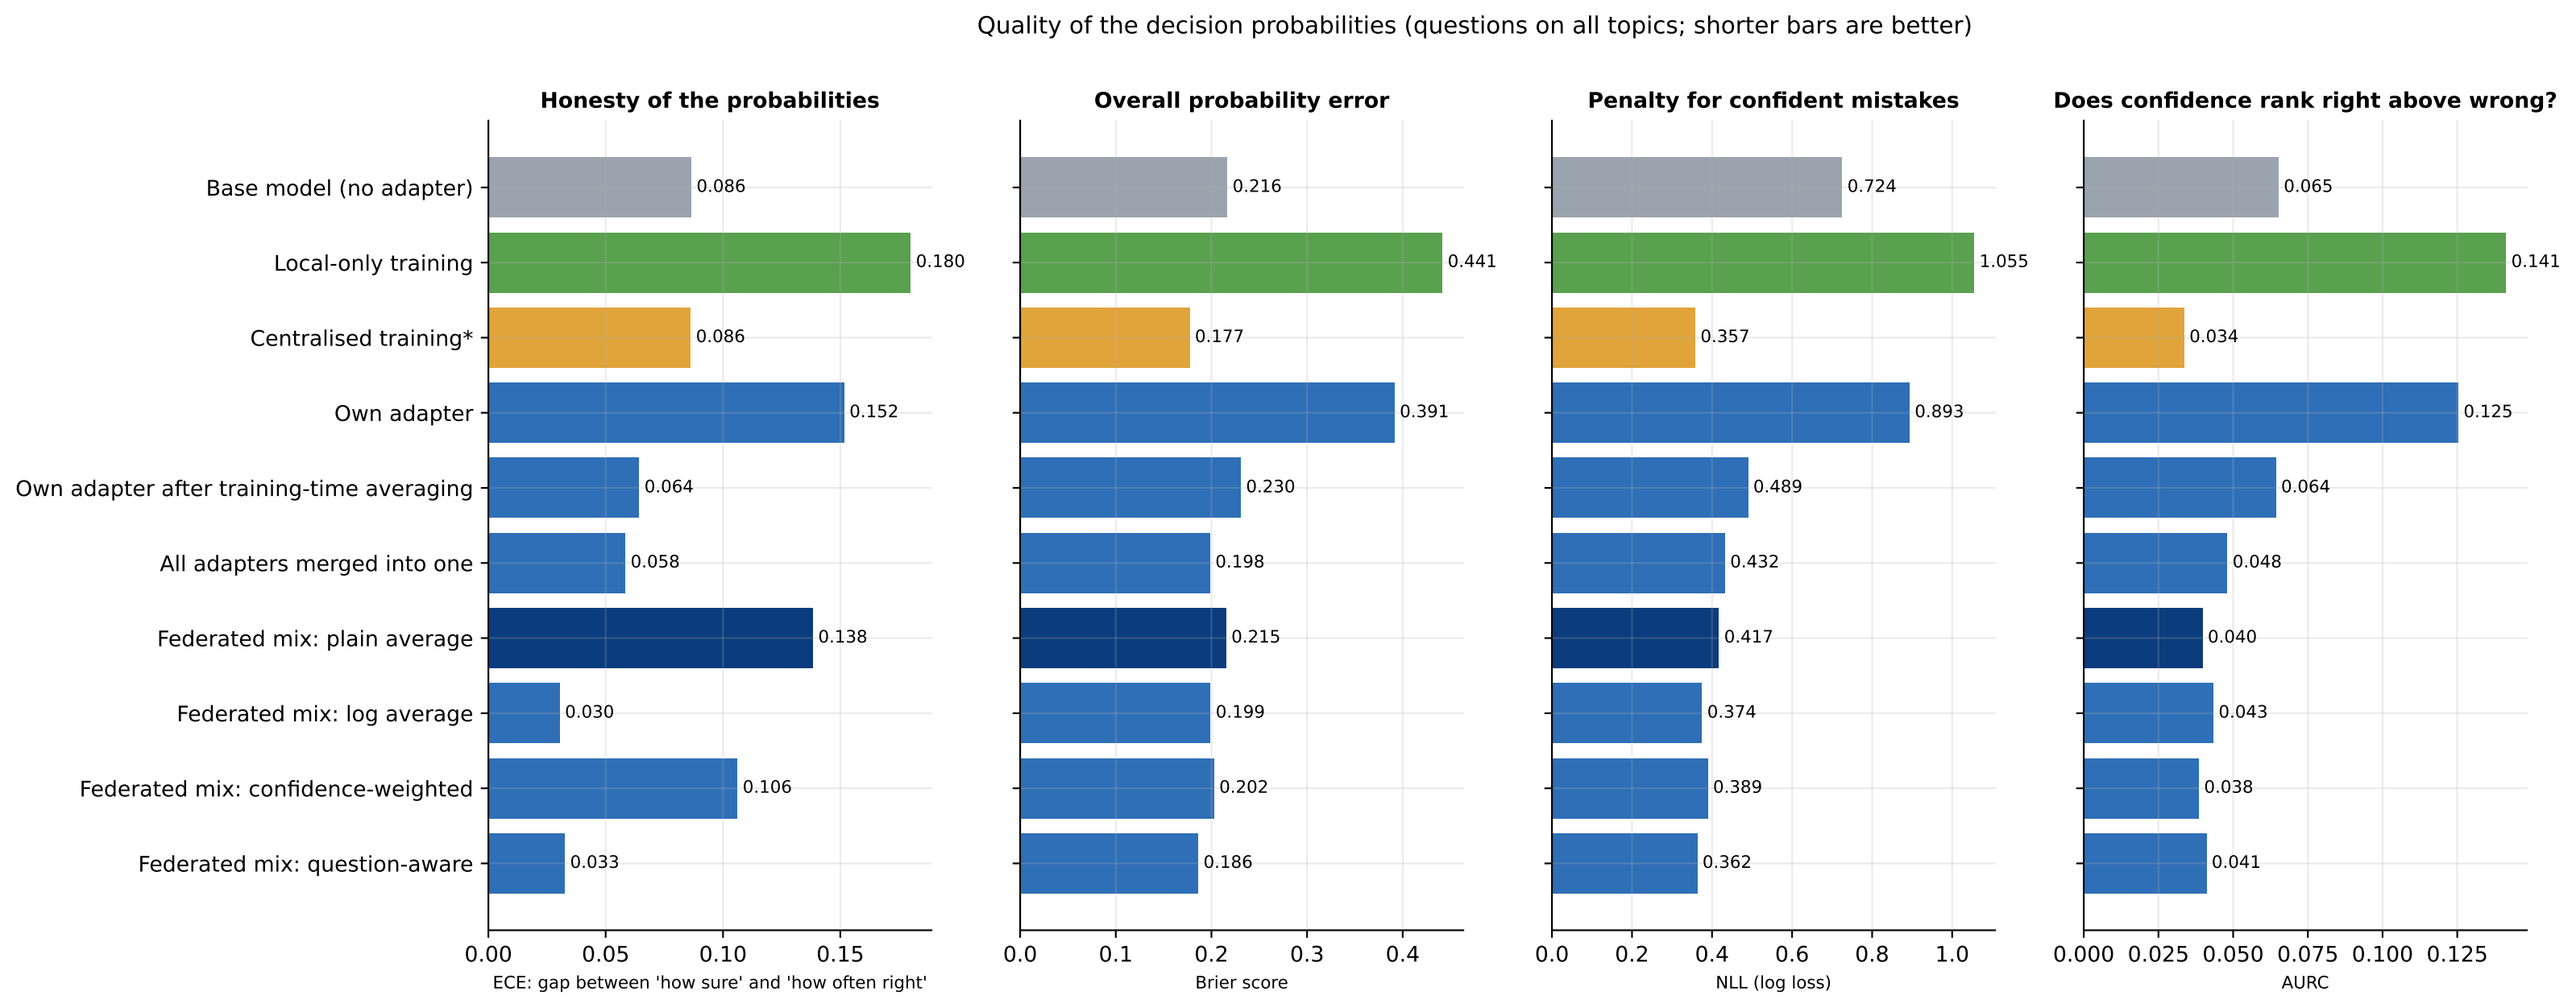
Quality of the decision probabilities (questions on all topics; shorter bars are better)
Honesty of the probabilities
Overall probability error
Penalty for confident mistakes
Does confidence rank right above wrong?
Base model (no adapter)
Local-only training
Centralised training*
Own adapter
Own adapter after training-time averaging
All adapters merged into one
Federated mix: plain average
Federated mix: log average
Federated mix: confidence-weighted
Federated mix: question-aware
0.086
0.180
0.086
0.152
0.064
0.058
0.138
0.030
0.106
0.033
0.216
0.441
0.177
0.391
0.230
0.198
0.215
0.199
0.202
0.186
0.724
1.055
0.357
0.893
0.489
0.432
0.417
0.374
0.389
0.362
0.065
0.141
0.034
0.125
0.064
0.048
0.040
0.043
0.038
0.041
0.00
0.05
0.10
0.15
ECE: gap between 'how sure' and 'how often right'
0.0
0.1
0.2
0.3
0.4
Brier score
0.0
0.2
0.4
0.6
0.8
1.0
NLL (log loss)
0.000
0.025
0.050
0.075
0.100
0.125
AURC