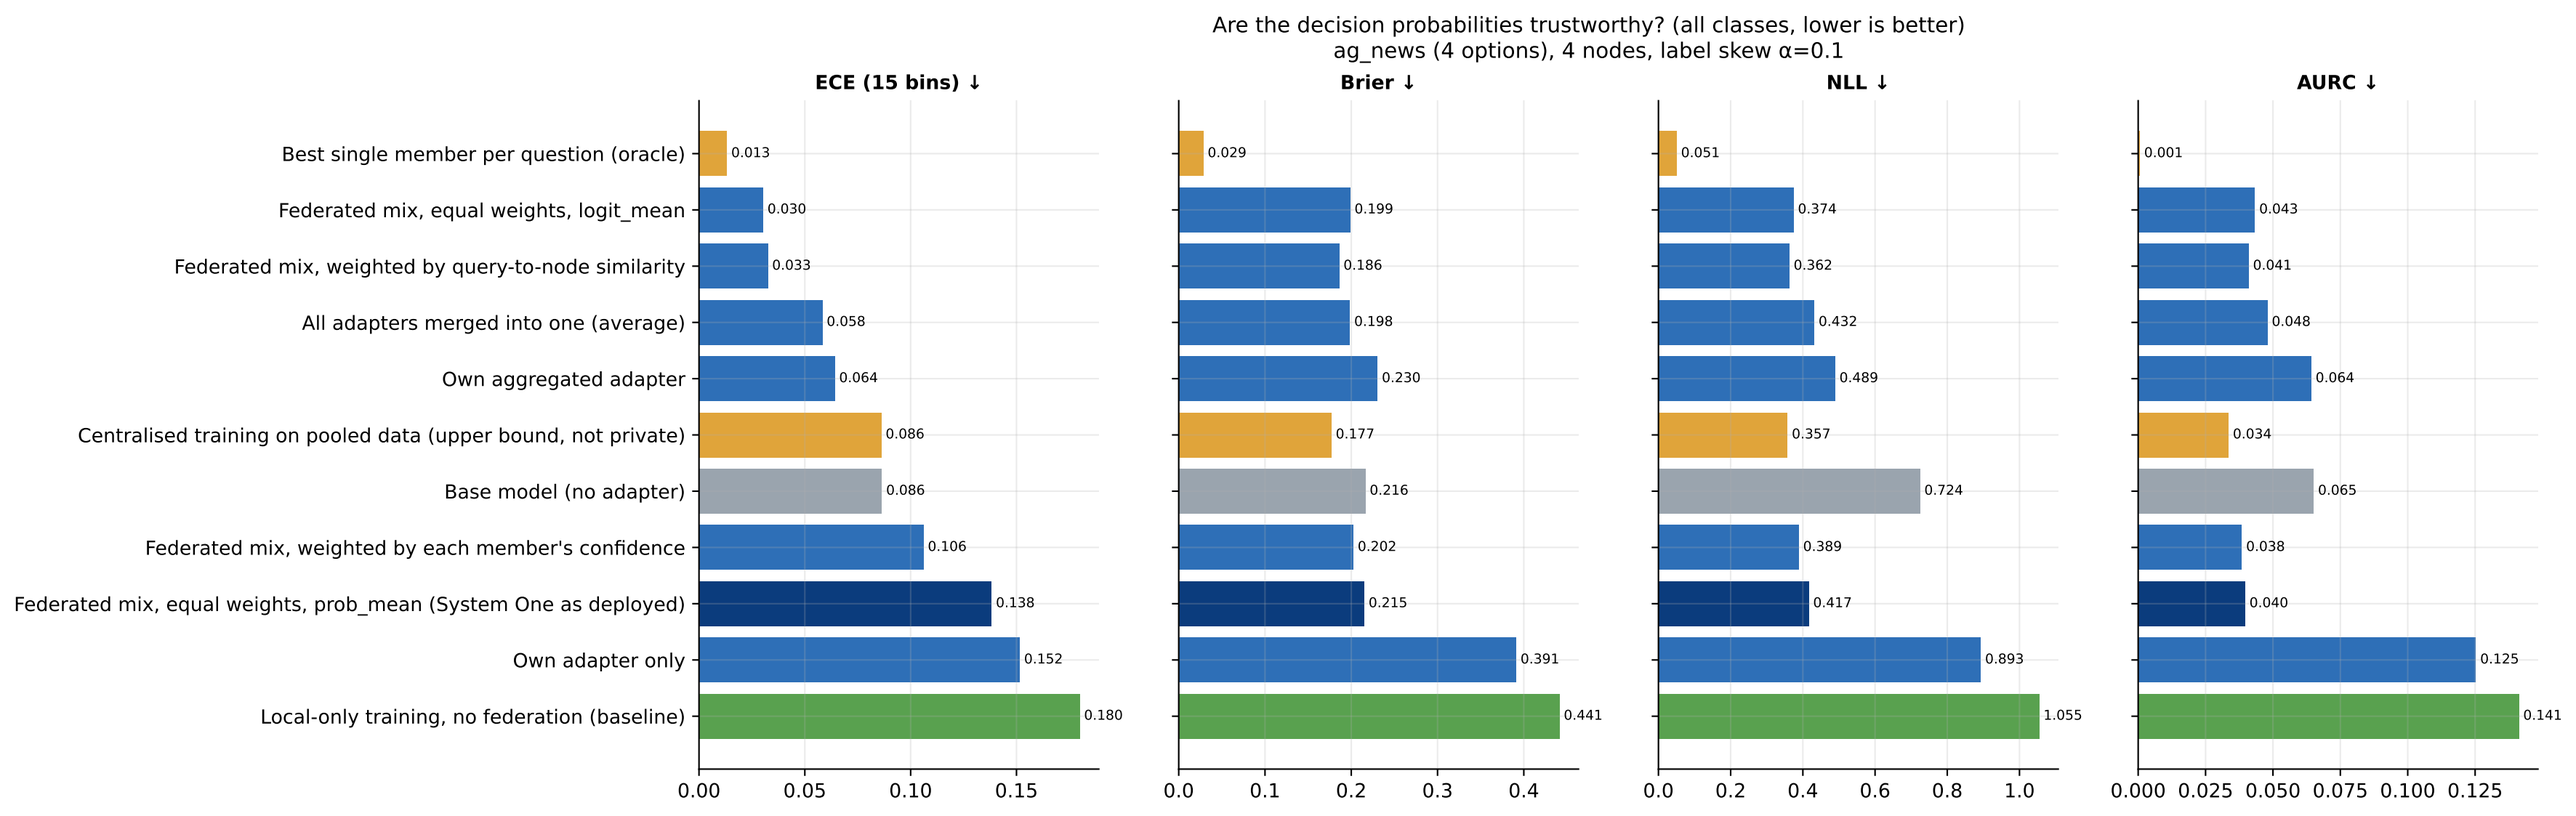
Are the decision probabilities trustworthy? (all classes, lower is better)
ag_news (4 options), 4 nodes, label skew α=0.1
ECE (15 bins) ↓
Brier ↓
NLL ↓
AURC ↓
Best single member per question (oracle)
Federated mix, equal weights, logit_mean
Federated mix, weighted by query-to-node similarity
All adapters merged into one (average)
Own aggregated adapter
Centralised training on pooled data (upper bound, not private)
Base model (no adapter)
Federated mix, weighted by each member's confidence
Federated mix, equal weights, prob_mean (System One as deployed)
Own adapter only
Local-only training, no federation (baseline)
0.013
0.030
0.033
0.058
0.064
0.086
0.086
0.106
0.138
0.152
0.180
0.029
0.199
0.186
0.198
0.230
0.177
0.216
0.202
0.215
0.391
0.441
0.051
0.374
0.362
0.432
0.489
0.357
0.724
0.389
0.417
0.893
1.055
0.001
0.043
0.041
0.048
0.064
0.034
0.065
0.038
0.040
0.125
0.141
0.00
0.05
0.10
0.15
0.0
0.1
0.2
0.3
0.4
0.0
0.2
0.4
0.6
0.8
1.0
0.000
0.025
0.050
0.075
0.100
0.125