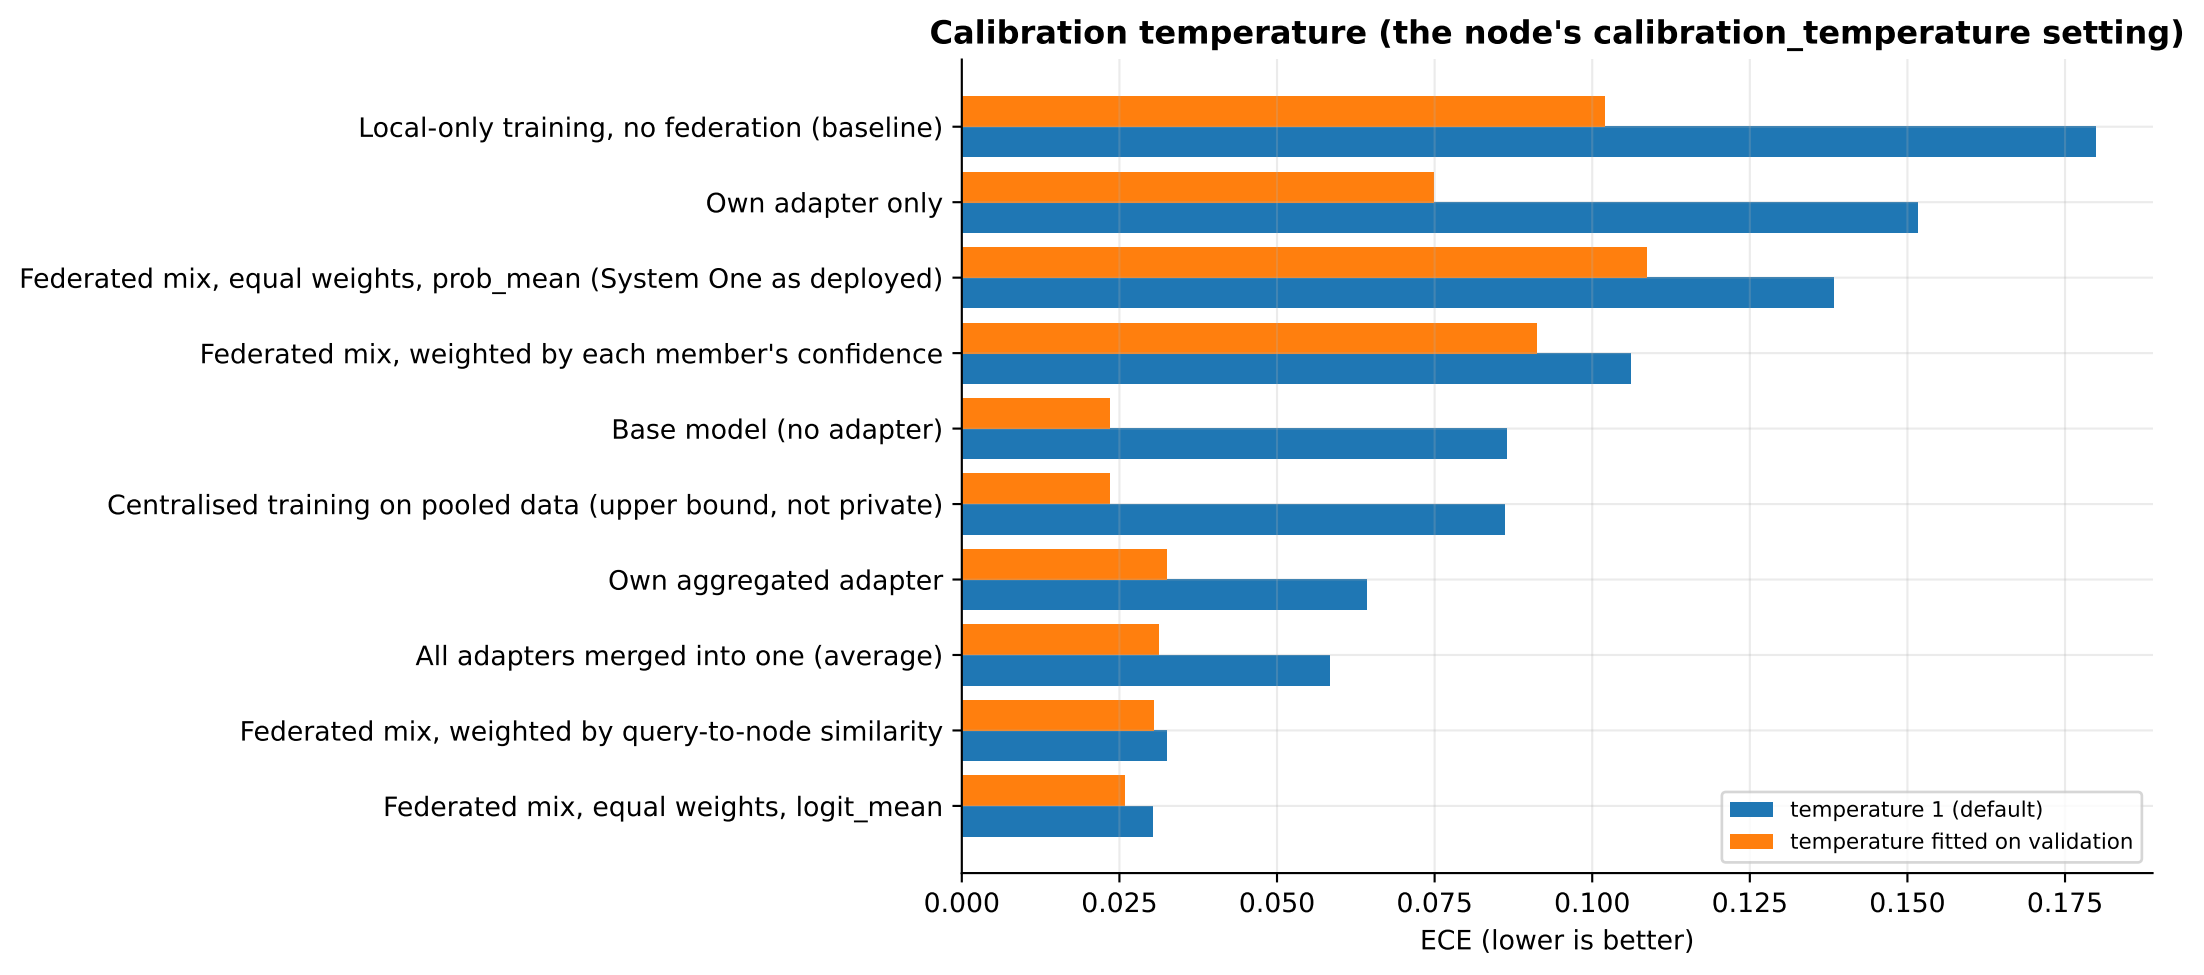
Calibration temperature (the node's calibration_temperature setting)
Local-only training, no federation (baseline)
Own adapter only
Federated mix, equal weights, prob_mean (System One as deployed)
Federated mix, weighted by each member's confidence
Base model (no adapter)
Centralised training on pooled data (upper bound, not private)
Own aggregated adapter
All adapters merged into one (average)
Federated mix, weighted by query-to-node similarity
Federated mix, equal weights, logit_mean
temperature 1 (default)
temperature fitted on validation
0.000
0.025
0.050
0.075
0.100
0.125
0.150
0.175
ECE (lower is better)