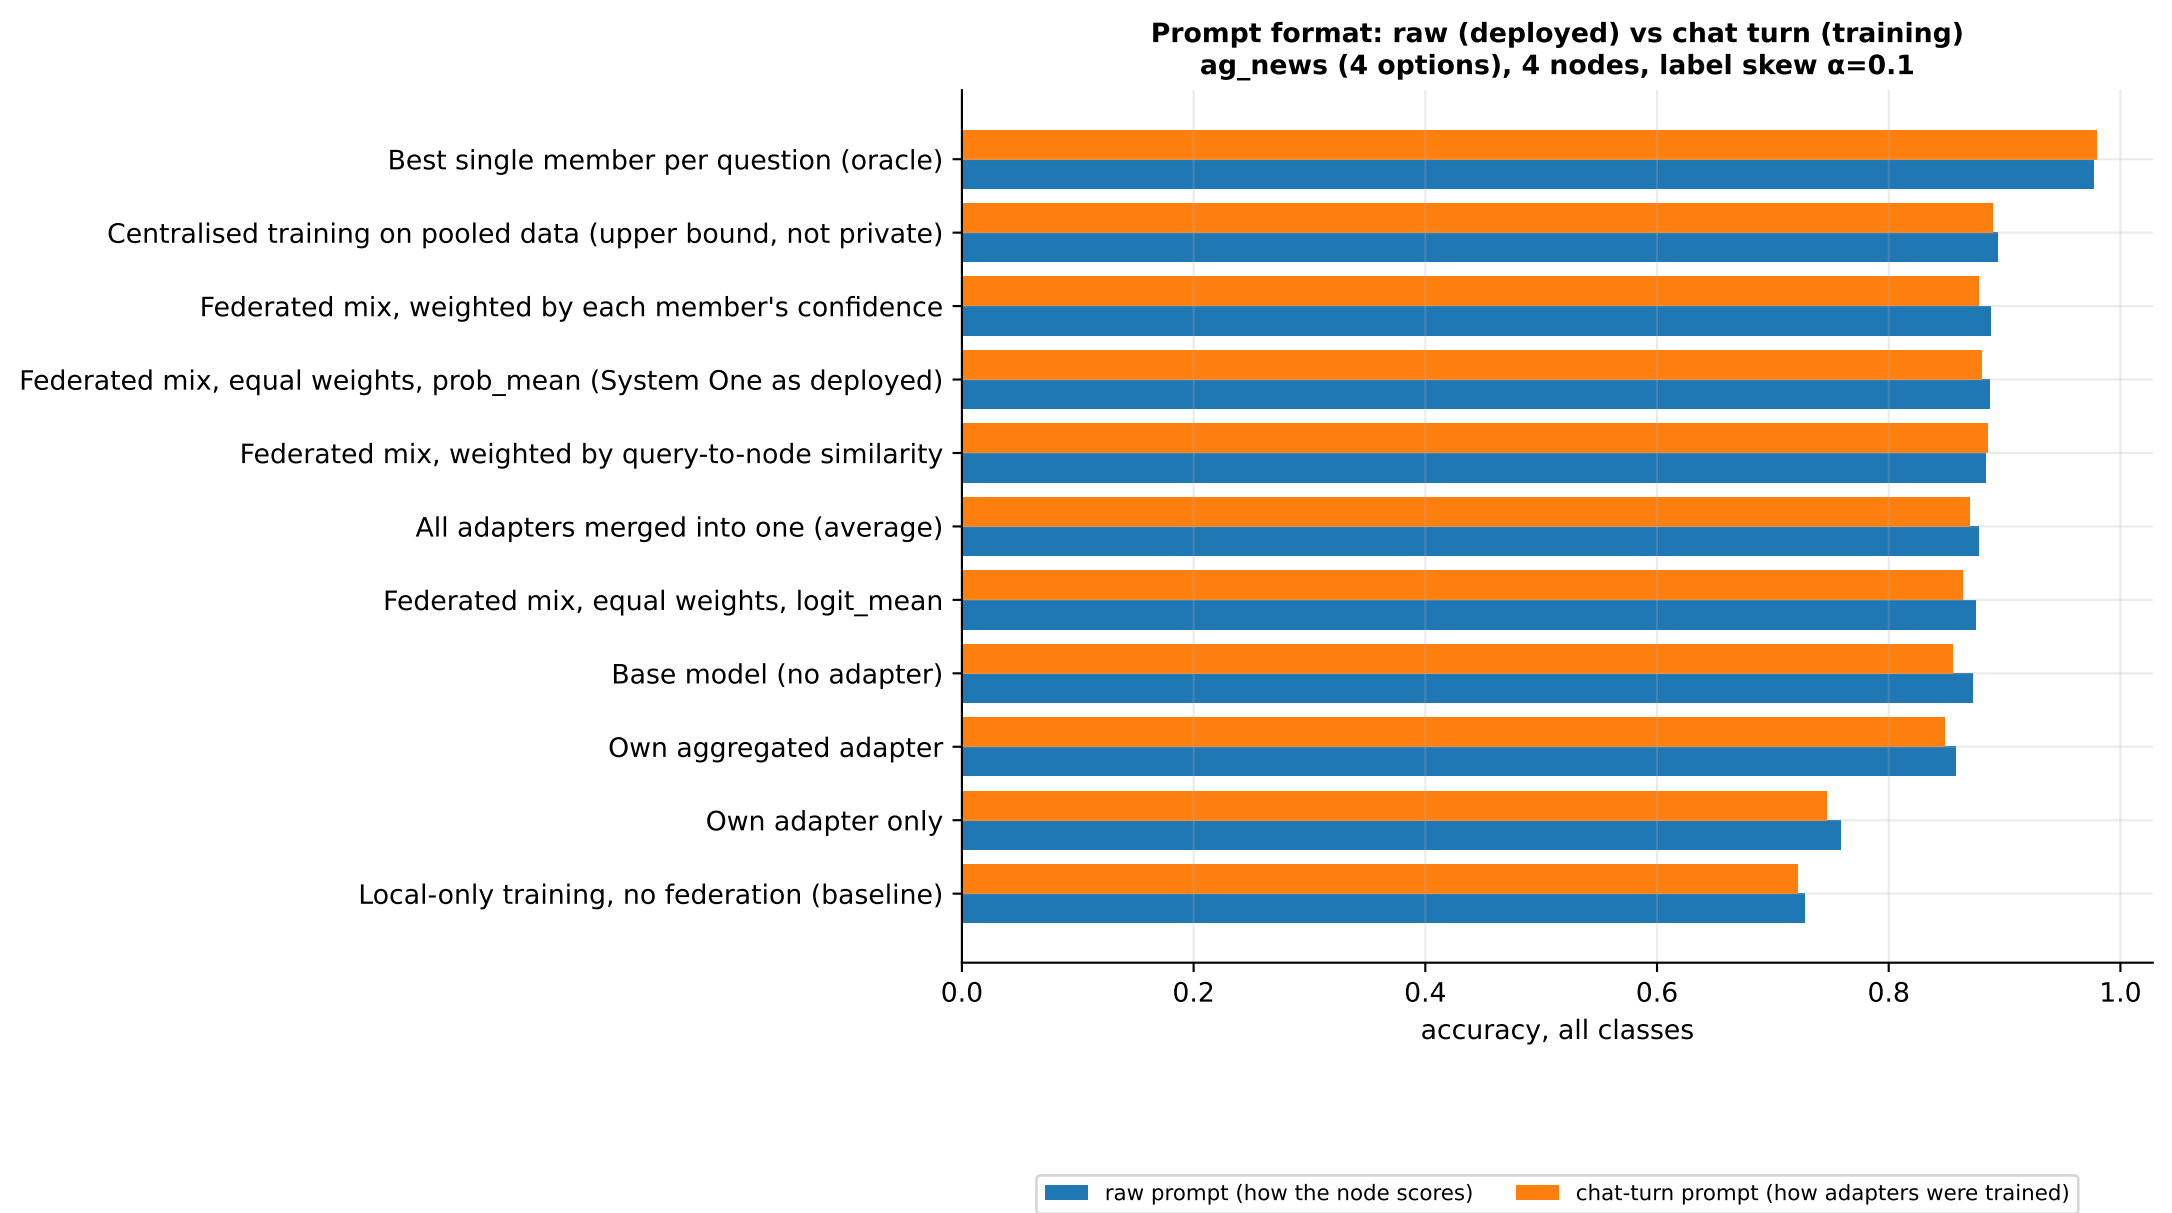
Prompt format: raw (deployed) vs chat turn (training)
ag_news (4 options), 4 nodes, label skew α=0.1
Best single member per question (oracle)
Centralised training on pooled data (upper bound, not private)
Federated mix, weighted by each member's confidence
Federated mix, equal weights, prob_mean (System One as deployed)
Federated mix, weighted by query-to-node similarity
All adapters merged into one (average)
Federated mix, equal weights, logit_mean
Base model (no adapter)
Own aggregated adapter
Own adapter only
Local-only training, no federation (baseline)
0.0
0.2
0.4
0.6
0.8
1.0
accuracy, all classes
raw prompt (how the node scores)
chat-turn prompt (how adapters were trained)