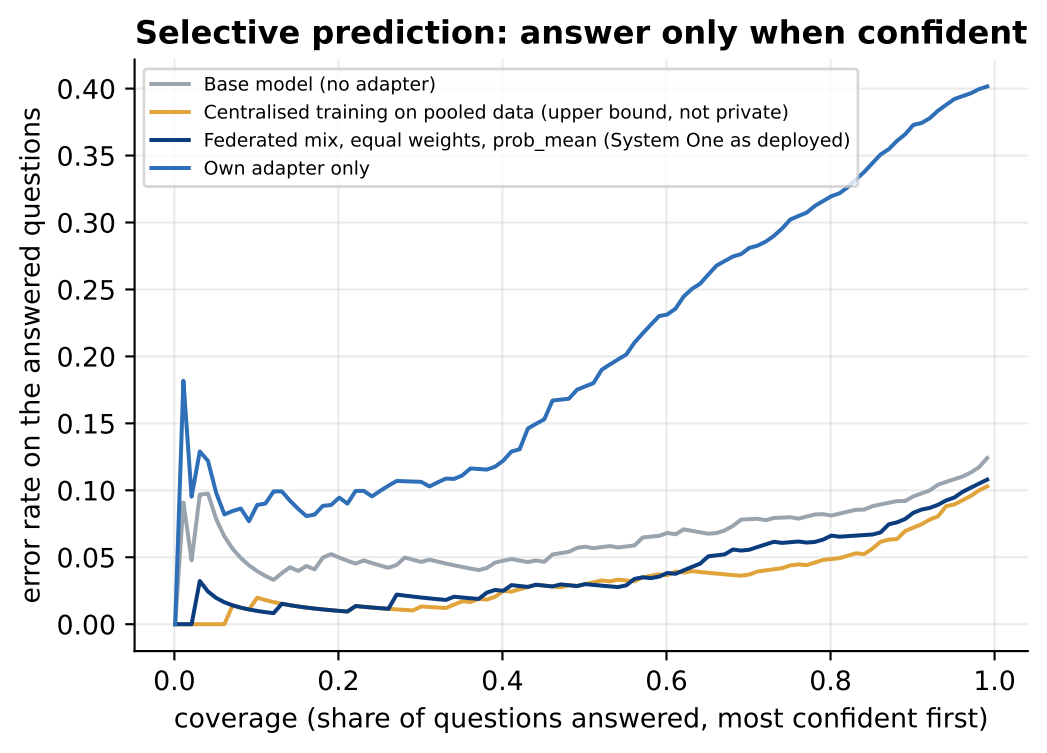
Selective prediction: answer only when confident
Base model (no adapter)
Centralised training on pooled data (upper bound, not private)
Federated mix, equal weights, prob_mean (System One as deployed)
Own adapter only
error rate on the answered questions
0.00
0.05
0.10
0.15
0.20
0.25
0.30
0.35
0.40
0.0
0.2
0.4
0.6
0.8
1.0
coverage (share of questions answered, most confident first)