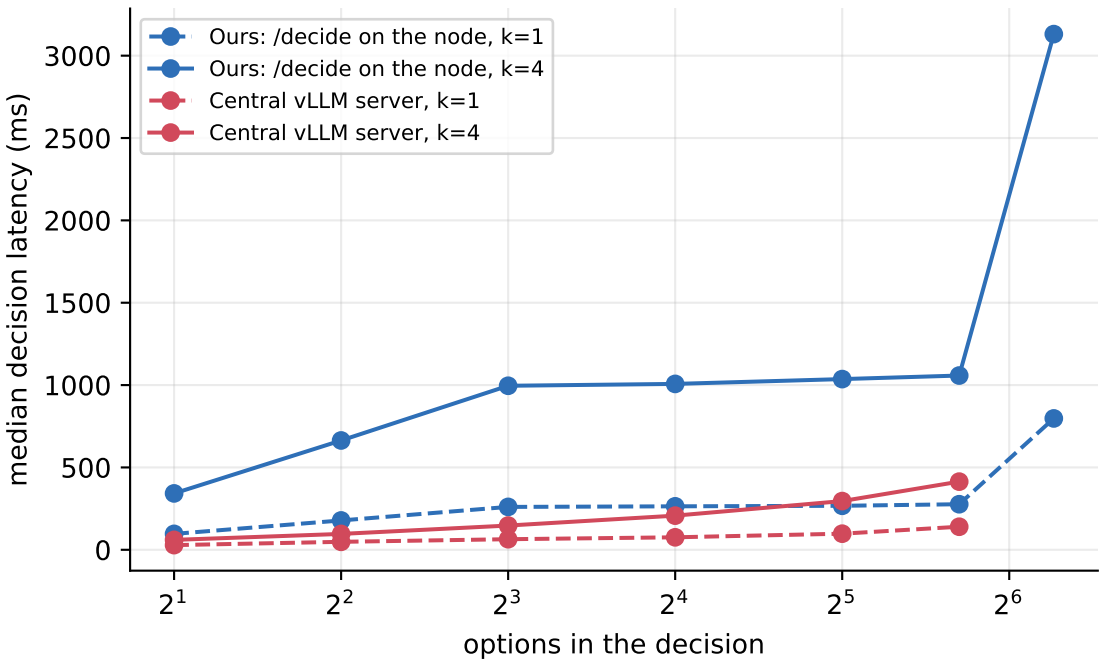
Latency vs number of options
Ours: /decide on the node, k=1
Ours: /decide on the node, k=4
Central vLLM server, k=1
Central vLLM server, k=4
median decision latency (ms)
0
500
1000
1500
2000
2500
3000
2^1
2^2
2^3
2^4
2^5
2^6
options in the decision