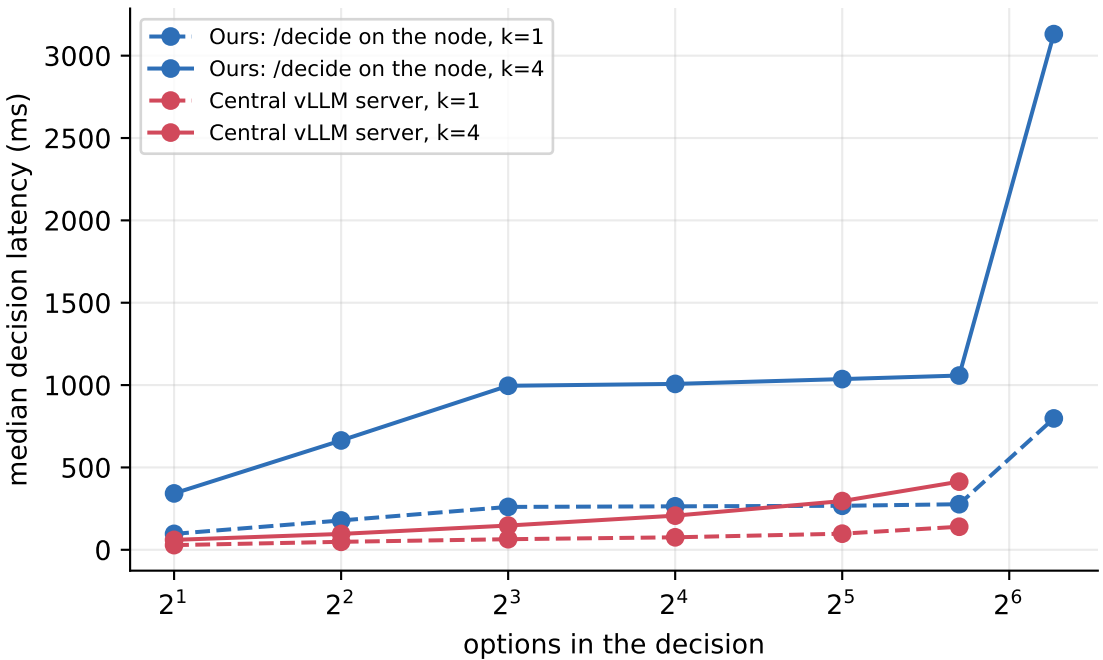
Latency vs number of options
Ours: /decide on the node, k=1
Ours: /decide on the node, k=4
Central vLLM server, k=1
Central vLLM server, k=4
median decision latency (ms)
0
500
1000
1500
2000
2500
3000
$2^1$
$2^2$
$2^3$
$2^4$
$2^5$
$2^6$
options in the decision