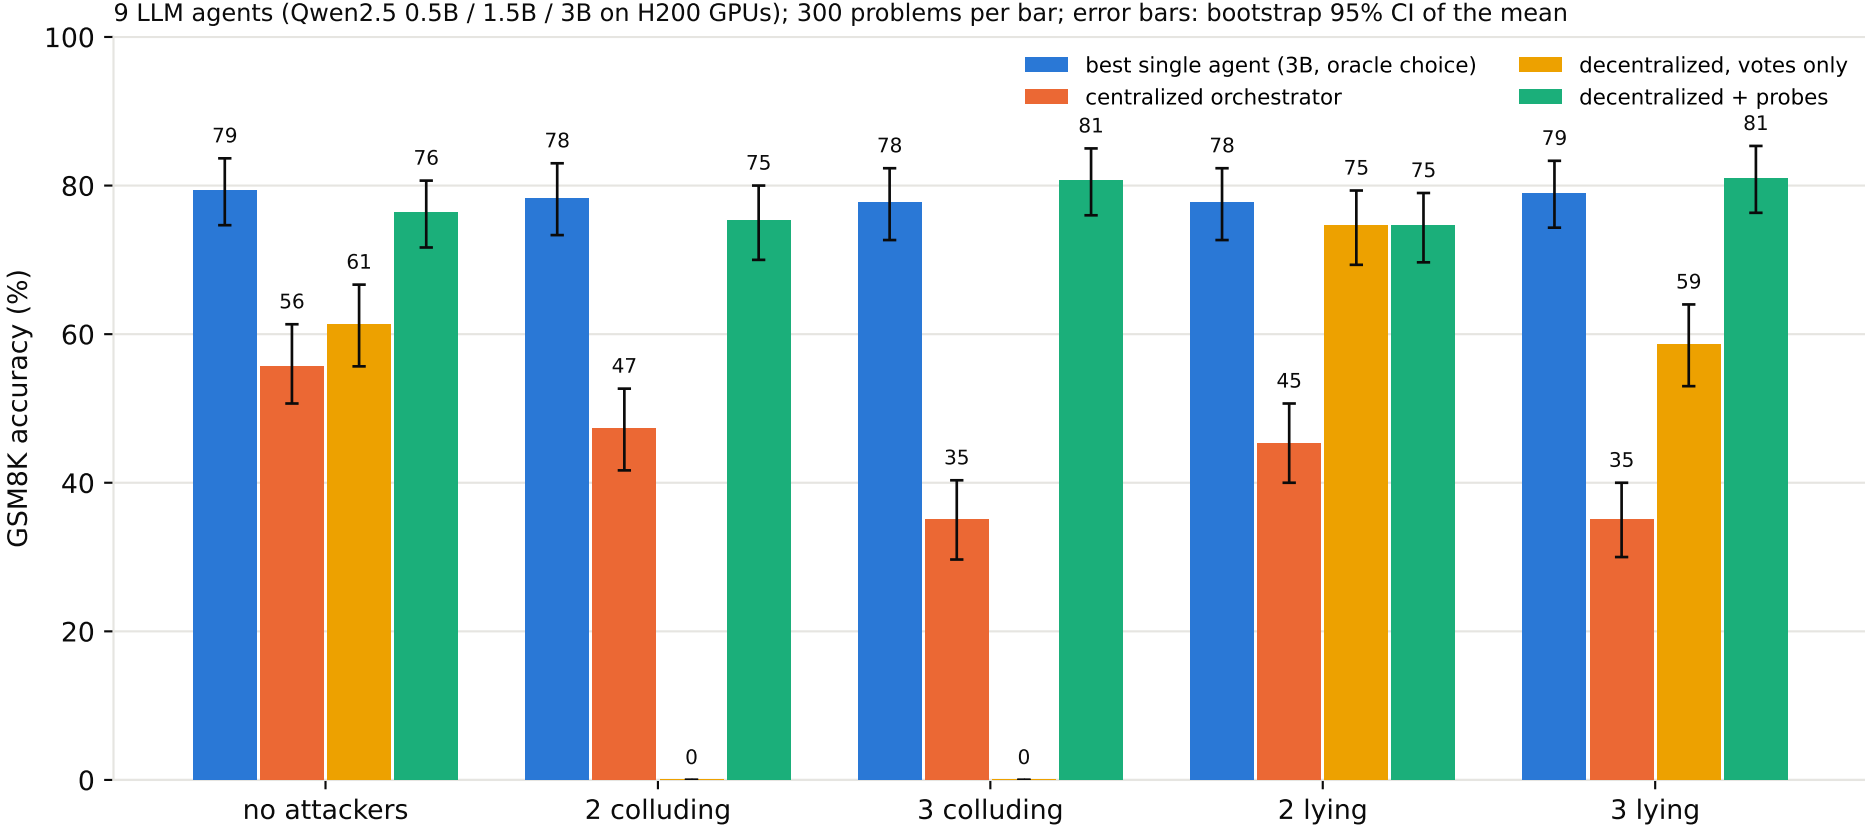
9 LLM agents (Qwen2.5 0.5B / 1.5B / 3B on H200 GPUs); 300 problems per bar; error bars: bootstrap 95% CI of the mean
best single agent (3B, oracle choice)
centralized orchestrator
decentralized, votes only
decentralized + probes
GSM8K accuracy (%)
100
80
60
40
20
0
79
56
61
76
78
47
0
75
78
35
0
81
78
45
75
75
79
35
59
81
no attackers
2 colluding
3 colluding
2 lying
3 lying